

1 Elementary sampling theory

Recall the basic rules:

- Product: $P(AB|C) = P(A|BC)P(B|C) = P(B|AC)P(A|C)$
- Sum: $P(AB) + P(\bar{A}|B) = 1$
- Extended sum: $P(A + B|C) = P(A|C) + P(B|C) - P(AB|C)$
- Principle of indifference: $P(H_i|B) = 1/N$, $1 \leq i \leq N$ - iff the set H_i is exhaustive and mutually exclusive
- Bernoulli urn rule: $P(A|B) = M/N$ - iff B specifies that A is true on some subset of H_i and false on the remaining $N - M$

1.1 Sampling without replacement

EXAMPLE: Bernoulli urn reexamined

- Propositions
 1. $B \equiv$ An urn contains N balls, identical except that they are labeled sequentially and M of them are coloured red, with the remaining $N - M$ coloured white. We draw a ball, observe and record its colour, and do not replace it. This is done until n balls are drawn, $0 \leq n \leq N$
 2. $R_i \equiv$ Red ball on the i th draw
 3. $W_i \equiv$ White ball on the i th draw
- Since only red or white can be drawn, we know that $P(R_i|B) + P(W_i|B) = 1$
- The propositions are related by the negations $\bar{R}_i = W_i$ and $R_i = \bar{W}_i$
- So, the first draw is then defined by the probabilities: $P(R_1|B) = M/N$ and $P(W_1|B) = 1 - M/N$
- Subsequent draws can be derived from the product rule: $P(R_1 R_2|B) = P(R_1|B)P(R_2|B)$, but need to take into account that the ball being drawn is not replaced.
 - Therefore, $P(R_1 R_2|B) = \frac{M}{N} \frac{M-1}{N-1}$
 - can be extended to r draws: $P(R_1 R_2 \dots R_r|B) = \frac{M!(N-r)!}{(M-r)!N!}$, where $r \leq M$
- What is the probability of drawing exactly r red balls in n draws, regardless of order?
 - must multiply by the binomial coefficient: $\binom{n}{r} = \frac{n!}{r!(n-r)!}$, which represents the number of possible orders of drawing r red balls in n draws, called the multiplicity of the event r .
 - * for example, to get three red in three draws, $\binom{3}{3} = 1$, can only happen in one way, $R_1 R_2 R_3$
 - * However, to get two red in three draws, $\binom{3}{2} = 3$, can happen in three ways, $R_1 R_2 W_3$, $R_1 W_2 R_3$, and $W_1 R_2 R_3$.
 - We can then derive an expression for drawing exactly R red balls in n draws, defined by the function $h(r|N, M, n) \equiv P(A|B)$

$$h(r|N, M, n) = \frac{\binom{M}{r} \binom{N-M}{n-r}}{\binom{N}{n}}$$
 - * This is called the hypergeometric distribution, often abbreviated as $h(r)$
 - it can be demonstrated that the hypergeometric distribution is symmetric on exchange of M and n , i.e. $h(r|N, M, n) = h(r|N, n, M)$. So, the probability of drawing ten ball balls from an urn containing 50 red ones is the same as drawing 50 balls from an urn containing 10 red ones.

- generalized hypergeometric distribution:

$$h(r_1 \dots r_k | N_1 \dots N_k) = \frac{\binom{N_1}{r_1} \dots \binom{N_k}{r_k}}{\binom{\sum N_i}{\sum r_i}}$$

where there are k different colours of N balls, drawn $n = \sum r_i$ times

- we can find the most probable value of r by setting $h(r') = h(r' - 1)$, and solving for r' . (this makes sense if you think about it like a distribution with a peak)
- the width of the distribution $h(r)$ gives an indication of the accuracy with which we can predict r
- Cumulative probability distribution: $H(R) \equiv \sum_{r=0}^R h(r)$, which is the probability of finding R or fewer red balls
 - $H(R)$ is a step function (think the first term in your first project)
- the median of the probability function is defined to be a number m which has equal probabilities associated with $(r < m)$ and $(r > m)$ (think a transition state)
- What is the probability of drawing a red ball on the second draw, $P(R_2|B)$, without knowledge of the first draw's result?
 - we know that either R_1 or W_1 is true, so $R_2 = (R_1 + W_1)R_2 = R_1R_2 + W_1R_2$
 - applying the product rule, we get

$$P(R_2|B) = P(R_1R_2|B) + P(W_1R_2|B) = P(R_2|R_1B)P(R_1|B) + P(R_2|W_1B)P(W_1|B)$$
 - but $P(R_2|R_1B) = \frac{M-1}{N-1}$ and $P(R_2|W_1B) = \frac{M}{N-1}$
 - so $P(R_2|B) = \frac{M-1}{N-1} \frac{M}{N} + \frac{M}{N-1} \frac{N-M}{N} = \frac{M}{N}$
 - notice that everything cancels out such that we have the same probability for red on the first and second draws.
 - this holds true generally, so the probability to draw red at any draw is the same, iff we do not know the result of any other draw

1.2 Logic vs Propensity

- since we know that knowledge of a earlier drawn ball's colour can change the probability of the current draw, can knowledge of a future ball's colour change the probability of the current ball?
- usually a fundamental axiom that future events can't change the probability of a current event
- but lets say we have an urn with two balls, one red and one white. the probability to draw a white ball with priors B is $P(W_1|B) = 1/2$. But what if we knew that the second draw was red? Then the probability becomes unity ($P(W_1|B) = 1$).
- So, while information about a later draw does not change the physical nature of the system, i.e. it does not change the number of balls of each colour in the system, it does change the state of knowledge of the system in the same way that knowledge of a past draw would
- this suggests that logical inference is fundamentally different from physical causation, e.g. physical influences propagate only forward in time while logical inferences propagate equally in either direction
- if the probability of an event is invariant under an permutation of the events (e.g. when the event took place), the probability distribution is called exchangeable. the hypergeometric distribution discussed above in the Bernoulli's urn example is exchangeable

1.3 Expectations

- definition: if a variable x can take on the set of values (x_1, \dots, x_n) , in n mutually exclusive and exhaustive situations, and we assign probabilities (p_1, \dots, p_n) , the quantity $\langle x \rangle = E(x) = \sum_{i=1}^n p_i x_i$ is the expectation value of x
 - (you know this from quantum mechanics)
 - it is a weighted average of the possible values of x , weighted by their corresponding probabilities
- this brings us to a an easier way to discuss probabilities with prior knowledge of an event at some later time not known
 - if $F = M/N$ of red balls is known, then $P(R_1|B) = F$
 - if F is unknown, $P(R_1|B) = \langle F \rangle$

1.4 Binomial Distribution

- the hypergeometric distribution takes into account the changing nature of the urn, i.e. drawing a ball changes the contents to $N-1$.
- but what if $N \gg n$? the probability changes very little, and in the limit $N \rightarrow \infty$, this becomes negligible
- thus the hypergeometric distribution simplifies to the binomial distribution (through a derivation i don't care to type):

$$h(r|N, M, n) \rightarrow b(r|n, f) \equiv \binom{n}{r} f^r (1-f)^{n-r} \text{ where } M/N \rightarrow f$$




Fig. 3.1. The hypergeometric distribution for $N = 15, 30, 100, \infty$.

Figure 1: Comparing the hypergeometric distribution to the binomial distribution (the hypergeometric distribution in the $N \rightarrow \infty$ limit)

- another limiting case, where $N_i \rightarrow \infty$, in such a way that $f_i \equiv \frac{N_i}{\sum N_j} = \text{constant}$,

$$m(r_1 \dots r_k | f_1 \dots f_k) = \frac{r!}{r_1! \dots r_k!} f_1^{r_1} \dots f_k^{r_k} \text{ where } r \equiv \sum r_i$$

- this is called the multinomial distribution

1.5 Sampling with replacement

- suppose instead of placing a drawn ball to the side, we instead put it back into the urn?
- with background information B' , the probability of drawing two red balls in succession is $P(R_1 R_2 | B') = P(R_1 | B') P(R_2 | R_1 B')$
 - clearly the first factor, $P(R_1 | B')$ is still M/N , but what about the second factor?
 - (there is an interesting digression on reality vs models wrt randomization, on pp. 73-75. randomization is not truly something found in nature, what we mean by randomization is that no human is able to implicitly or explicitly influence the results)
 - if we assume the urn to be randomized upon replacement, information of R_1 is irrelevant to draw 2, so $P(R_2 | R_1 B') = P(R_2 | B') = M/N$. this holds true generally for any draw.
 - the probability for drawing exactly r balls in n trials is simply $\binom{n}{r} \left(\frac{M}{N}\right)^r \left(\frac{N-M}{N}\right)^{n-r}$, which is just the binomial distribution
 - thus, the probability of drawing r red balls with replacement is the same as drawing r red balls without replacement in the limit $N \rightarrow \infty$
 - this approximation, however, can accumulate errors for large n
 - * suppose that drawing and replacing a red ball increases the probability of drawing a red ball on the next draw by some small amount $\epsilon > 0$, while drawing a white ball decreases it by some small amount $\delta > 0$ (think of this as nonoptimal shaking of the urn)
 - * letting C be the background information described above, we have the following probabilities:

$$\begin{aligned} P(R_k | R_{k-1} C) &= p + \epsilon & P(R_k | W_{k-1} C) &= p - \delta \\ P(W_k | R_{k-1} C) &= 1 - p - \epsilon & P(W_k | W_{k-1} C) &= 1 - p + \delta \end{aligned}$$

where $p \equiv M/N$ (this is referenced below)

- * from this, the probability of drawing r red balls and $(n - r)$ white balls in any order is $p(p + \epsilon)^c (p - \delta)^{c'} (1 - p + \delta)^w (1 - p - \epsilon)^{w'}$, where c is the number of red draws preceded by red ones, c' is the number of red preceded by white, w is the number of white preceded by white, and w' is the number of white preceded by red.
- * we can additionally see that $c + c' = \lfloor \frac{n-r}{r} \rfloor$ and $w + w' = \lfloor \frac{n-r}{n-r-1} \rfloor$, where the upper and lower cases are when the first draw is red or white, respectively.
 - when r and $(n - r)$ are small, ϵ and δ are negligible, and the expression simplifies to $p^r (1 - p)^{n-r}$, as in the binomial distribution above
 - but as $r, n \rightarrow \infty$, we can use the relation $\left(1 + \frac{\epsilon}{p}\right)^c \approx \exp\left(\frac{\epsilon c}{p}\right)$, so that the probability goes to $p^r (1 - p)^{n-r} \exp\left(\frac{\epsilon c - \delta c'}{p} + \frac{\delta w - \epsilon w'}{1-p}\right)$
 - thus, the probability now depends on the order, and depending on ϵ and δ , the deviation from the binomial distribution can be large
- Let's see how this affects previous calculations
 - * for the first draw we still have $p = P(R_1 | C) = M/N$ and $q = 1 - p = P(W_1 | C) = \frac{N-M}{N}$
 - * but for the second trial we have $P(R_2 | C) = p + (p\epsilon - q\delta)$, and $P(R_3 | C) = p + (1 + \epsilon + \delta)(p\epsilon - q\delta)$ for the third. does $P(R_k | C)$ approach some limit as $k \rightarrow \infty$?
 - * if we write the probabilities for the k th trial as a vector $V_k \equiv \begin{bmatrix} P(R_k | C) \\ P(W_k | C) \end{bmatrix}$, the set of the four probabilities above can be written in matrix form $V_k = M V_{k-1}$ where $M = \begin{bmatrix} p + \epsilon & p - \delta \\ q - \epsilon & q + \delta \end{bmatrix}$
 - This is a Markov chain of probabilities (!), and M is called the transition matrix.
 - So draws after the first can be expressed like $V_k = M^{k-1} V_1$.
 - To find a general solution, we need to find the eigenvectors and eigenvalues of M : $C(\lambda) \equiv \det(M_{ij} - \lambda \delta_{ij}) = \lambda^2 - \lambda(1 + \epsilon + \delta) + (\epsilon + \delta)$, so $\lambda_1 = 1$ and $\lambda_2 = \epsilon + \delta$

- the eigenvectors are $x_1 = \begin{pmatrix} p-\delta \\ q-\epsilon \end{pmatrix}$ and $x_2 = \begin{pmatrix} 1 \\ -1 \end{pmatrix}$, which are not orthogonal.
- using these to define the transition matrix, $S = \begin{pmatrix} p-\delta & 1 \\ q-\epsilon & -1 \end{pmatrix}$, we can diagonalize M , $(S^{-1}MS)$, to eventually get to the general solution:

$$P(R_k|C) = \frac{(p-\delta) - (\epsilon+\delta)^{k-1}(p\epsilon - q\delta)}{1 - \epsilon - \delta}$$

- when $\epsilon = \delta = 0$, this simplifies to $P(R_k|C) = p$
- * interesting to note that $\epsilon + \delta = 1$ iff $\epsilon = q$ and $\delta = p$ (since $0 \leq p \leq 1$, ϵ and δ must be bounded to the interval $[-1,1]$), the transition matrix simplifies to $M = \begin{bmatrix} p+\epsilon & p-\delta \\ q-\epsilon & q+\delta \end{bmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$, and no transitions occur
- * likewise, if $\epsilon + \delta = -1$, $M = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$, and nothing but transitions occur, i.e. colours alternate after the first draw
- * in the realistic case where $0 < |\epsilon + \delta| < 1$, the general solution attenuates exponentially with k to give the limit $P(R_k|C) \rightarrow \frac{p-\delta}{1-\epsilon-\delta}$
 - it is clear that this limiting distribution is not exchangeable, since the conditional probabilities depend on the separation $|k-j|$ of each draw.
 - let us consider $P(R_k|R_jC)$. the general solution is

$$P(R_k|R_jC) = \frac{(p-\delta) + (\epsilon+\delta)^{k-j}(q-\epsilon)}{1 - \epsilon - \delta}$$

where $j < k$

- this approaches the limit above. since we have seen that $P(R_k|C) \neq P(R_j|C)$, it follows that $P(R_j|R_kC) \neq P(R_k|R_jC)$, due to the product rule.
- this is the “baby” form of the irreversible Markov process, where the current state depends only on the previous state and is asymmetric around the current state.